

CAB FARE PREDICTION

Shridhar S

11 September 2019

INDEX

1	Introduction	3
1.1	Problem statement	3
1.2	Data	3
2	Methodology	5
2.1	Pre Processing	5
2.2	Missing Value Analysis	5
2.3	Feature Engineering	5
2.4	Outliers	6
2.5	Distribution of Continuous Variables	8
2.6	Distribution of Categorical Variables	9
2.7	Feature Selection	10
3	Modelling	12
3.1	Model Selection	12
3.2	Multiple Linear Regression	12
3.3	Random Forest	13
4	Conclusion	14
4.1	Model Evaluation	14
4.2	Error Metrics	14
4.3	Mean Absolute Percentage Error(MAPE)	15
4.4	Model Selection	15
5	R Code	16
6	Python Code	37

Introduction

1.1 Problem Statement :

Predicting a Cab fare for any trip before the users book is a crucial step for any cab rental service. It gives much clarity and cost estimation to the users about the trip they are going to take. We are a Start-up Cab company. Having successfully run our Pilot project across the country, we have the same data at our disposal using which, we need to come up with a model to predict cab fare for each trip

1.2 Data

Our aim is to build a model which will predict the cab fare for each trip. We have a sample dataset given below to have an idea of how our dataset looks and what are the variables that help us in predicting the cab fare.

Table 1.1 : Cab fare prediction sample data

fare_amount	pickup_datetime	pickup_longitude	pickup_latitude	dropoff_longitude	dropoff_latitude	passenger_count
4.5	2009-06-15 17:26:21 UTC	-73.844311	40.721319	-73.84161	40.712278	1
16.9	2010-01-05 16:52:16 UTC	-74.016048	40.711303	-73.979268	40.782004	1
5.7	2011-08-18 00:35:00 UTC	-73.982738	40.76127	-73.991242	40.750562	2
7.7	2012-04-21 04:30:42 UTC	-73.98713	40.733143	-73.991567	40.758092	1
5.3	2010-03-09 07:51:00 UTC	-73.968095	40.768008	-73.956655	40.783762	1
12.1	2011-01-06 09:50:45 UTC	-74.000964	40.73163	-73.972892	40.758233	1
7.5	2012-11-20 20:35:00 UTC	-73.980002	40.751662	-73.973802	40.764842	1
16.5	2012-01-04 17:22:00 UTC	-73.9513	40.774138	-73.990095	40.751048	1

As we can see from the sample data, we have a total of 7 variables in the dataset.

Table 1.2 Unique Variables

1	fare_amount
2	pickup_datetime
3	pickup_longitude
4	pickup_latitude
5	dropoff_longitude
6	dropoff_latitude
7	passenger_count

Following is the explanation of the variables.

pickup_datetime - timestamp value indicating when the cab ride started.

pickup_longitude - float for longitude coordinate of where the cab ride started.

pickup_latitude - float for latitude coordinate of where the cab ride started.

dropoff_longitude - float for longitude coordinate of where the cab ride ended.

dropoff_latitude - float for latitude coordinate of where the cab ride ended.

passenger_count - an integer indicating the number of passengers in the cab ride.

Methodology

2.1 Pre – Processing

The data we receive would not be clean or in-line with what we need for analysis. In order to make the data best for visualization, graphs and plots, we need to clean the data and make necessary changes. This is called Data Pre-Processing and it is one of the most important steps in Data Science. We have to check for missing values, outliers and treat them to keep the sanity of the data. Also we need to extract new features out of given data and see how they can help us in predicting the dependent variables.

2.2 Missing Value Analysis

We see that we have 25 missing values(NA's) in fare_amount and 55 missing values in passenger_count. Usually we deal with missing values by substituting them with mean, median or nearest value. Since in our analysis, fare_amount is the dependent variable and also the missing values are not much in number, we will delete the rows with missing values instead.

2.3 Feature Engineering

a) We can extract some features like year, month, day, hour, date out of the timestamp provided in the pickup_datetime variable given in the data. This gives us some valuable inputs as to how the fare_amount depends on different part of the day, different months of the year, on weekdays and weekends etc.

b) We can also extract distance between two geocodes using Haversine formula.

Haversine formula to find distance between two points on a sphere

The **Haversine** formula calculates the shortest distance between two points on a sphere using their latitudes and longitudes measured along the surface. It is important for use in navigation. The haversine can be expressed in trigonometric function as:

$$\text{haversine}(\theta) = \sin^2\left(\frac{\theta}{2}\right)$$

The haversine of the central angle (which is d/r) is calculated by the following formula:

$$\left(\frac{d}{r}\right) = \text{haversine}(\Phi_2 - \Phi_1) + \cos(\Phi_1)\cos(\Phi_2)\text{haversine}(\lambda_2 - \lambda_1)$$

where r is the radius of earth(6371 km), d is the distance between two points, ϕ_1, ϕ_2 is latitude of the two points and λ_1, λ_2 is longitude of the two points respectively.

Solving d by applying the inverse haversine or by using the inverse sine function, we get:

$$d = r \text{hav}^{-1}(h) = 2r \sin^{-1}(\sqrt{h})$$

or

$$d = 2r \sin^{-1}\left(\sqrt{\sin^2\left(\frac{\Phi_2 - \Phi_1}{2}\right) + \cos(\Phi_1)\cos(\Phi_2)\sin^2\left(\frac{\lambda_2 - \lambda_1}{2}\right)}\right)$$

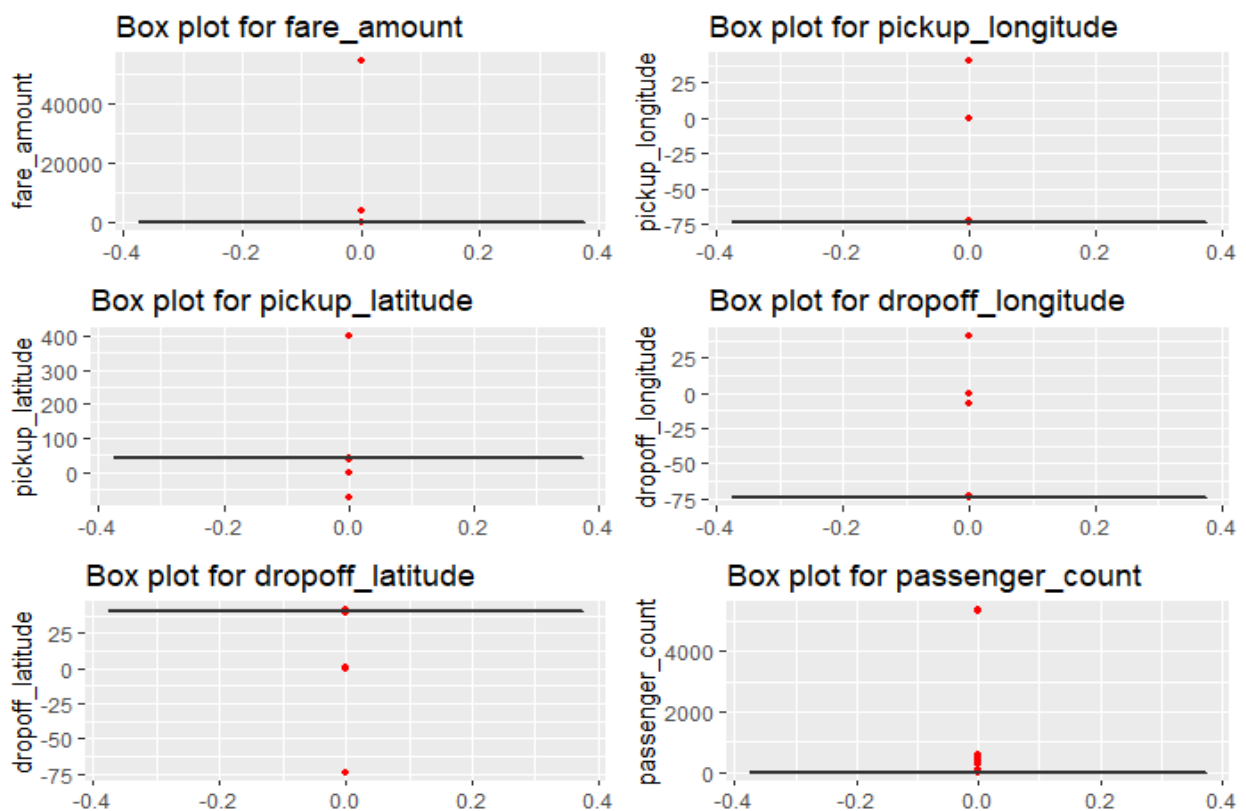
2.4 Outlier Analysis

Outliers are those values which lie on the outside of the given range. These extreme values deviate us from the actual observations and the actual pattern of the observations. It's important to treat the outliers as this affects our prediction adversely in the prediction stages. Outliers can be easily identified with the help of statistical methods and graphs(Boxplots). Boxplots is one such graphs which contains 3 quantiles(Q1, Median, Q3). Q1 is the 25th percentile of the variable and Q3 is the 75th percentile. The difference between Q3 and Q1 is called the Inter Quantile Range or IQR. So the value which is outside of Q3 by 1.5 times the IQR is called the Upper limit. Values which are outside of Q1 by 1.5 times the IQR is called Lower limit.

Outliers for continuous variables for using statistical mean, median, max and min.

```
> summary(numeric_data)
fare_amount      pickup_longitude pickup_latitude dropoff_longitude dropoff_latitude passenger_count
Min.   :   -3.00   Min.   : -74.44   Min.   : -74.01   Min.   : -74.43   Min.   : -74.01   Min.   :   0.000
1st Qu.:    6.00   1st Qu.: -73.99   1st Qu.:  40.73   1st Qu.: -73.99   1st Qu.:  40.73   1st Qu.:   1.000
Median :    8.50   Median : -73.98   Median :  40.75   Median : -73.98   Median :  40.75   Median :   1.000
Mean   :   15.03   Mean   : -72.46   Mean   :  39.92   Mean   : -72.46   Mean   :  39.90   Mean   :   2.623
3rd Qu.:   12.50   3rd Qu.: -73.97   3rd Qu.:  40.77   3rd Qu.: -73.96   3rd Qu.:  40.77   3rd Qu.:   2.000
Max.   : 54343.00   Max.   :  40.77   Max.   : 401.08   Max.   :  40.80   Max.   :  41.37   Max.   : 5345.000
> |
```

Outlier for the continuous variables in the data using boxplot.



Once the values which are Once the outliers are identified, we can treat them by replacing them with suitable values such as upper limit or lower limit of the variables.

2.5 Distribution of Continuous Variables

From the below graphs we can see the distribution of continuous variables like fare_amount, pickup_longitude, pickup_latitude etc.

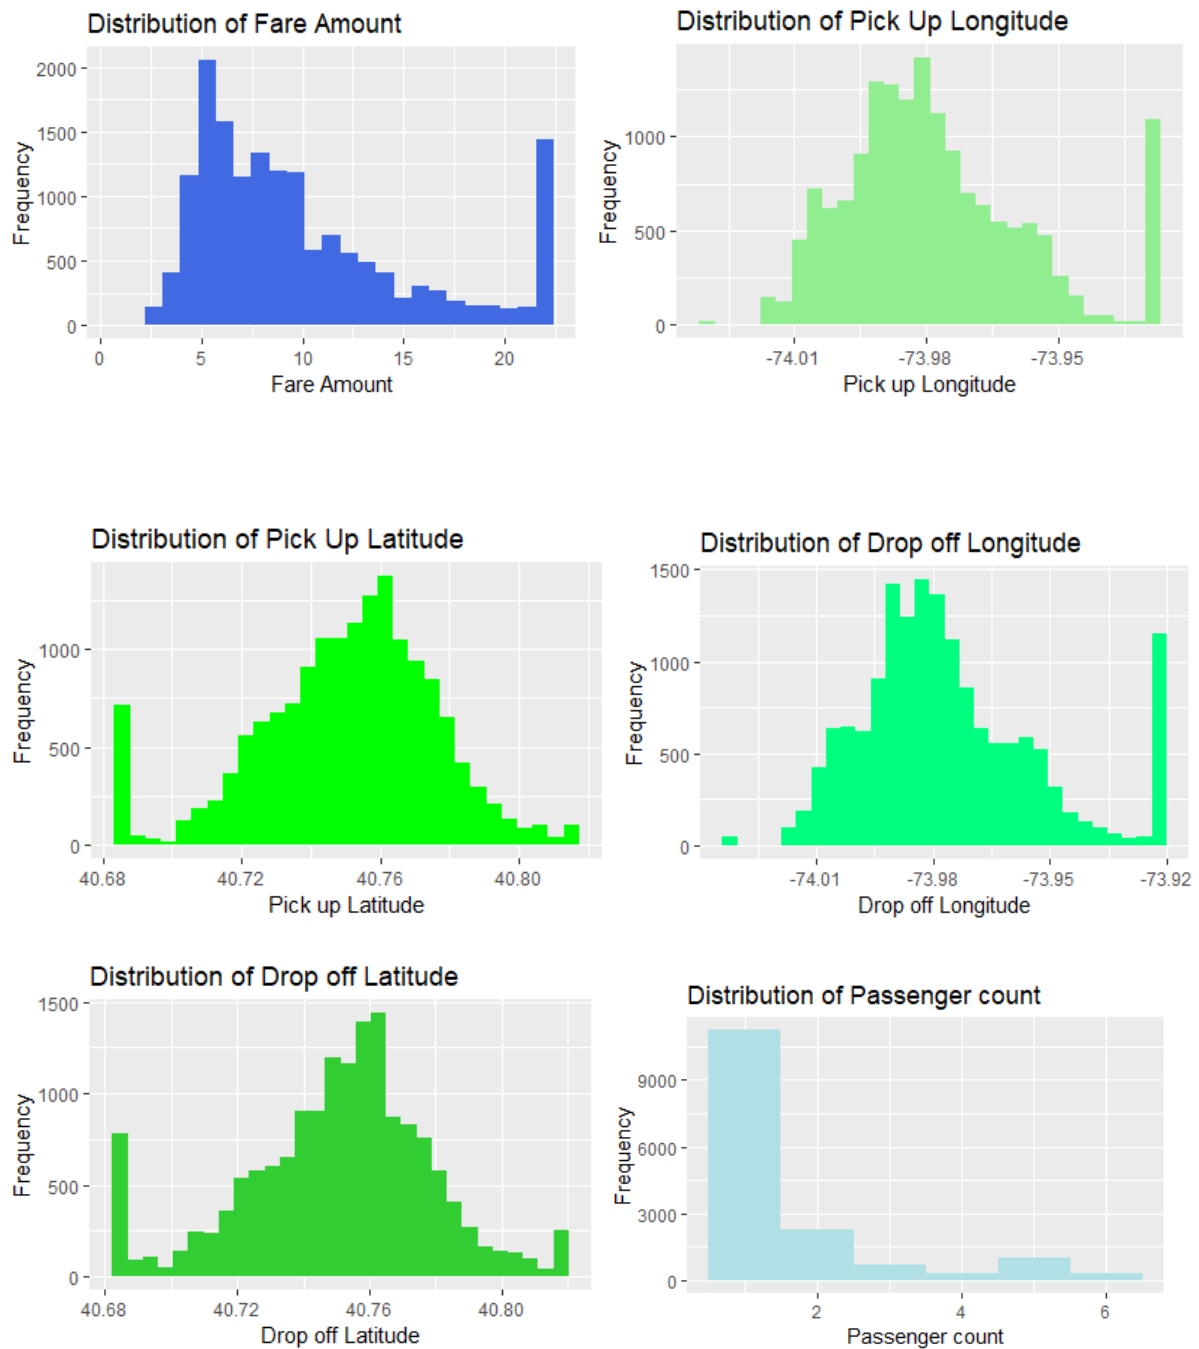




Fig. 2.1 : Categorical variables vs count



We can see the frequency distribution of all the continuous variables in the above graph. We observe that all the variables are normally distributed (more or less) and that is as expected except for the distance. We can see that Distance is right skewed and this skewness can be removed by log transformation. Also the Passenger count distribution need not be a normal distribution as there could be a possibility that a single passenger took most of the trips compared to the other numbers

2.6 Distribution of Categorical Variables

In the below graph we can see the distribution of continuous variables such as year, month, day, date, hour.






Fig. 2.2 : Categorical variables vs fare amount

We see that the total fare amount increased over the years, however by 2014-15 it reduced. We can see less fare amount in July and August. Fare amount remained same throughout the week. When it comes to the hours, we can see lesser fare amount in the early morning and good number of booking after the evening.

2.7 Feature Selection

As we looked the dependency of dependent variable on other variable, we could see what are the variables which are affecting the number of counts. It is important for us to select the variables which are highly related to the dependent variable and reduce any error that we may have. This can be done using the Correlation plot or Correlograms.

```
[1] "date"
      Df Sum Sq Mean Sq F value Pr(>F)
train[, i]      30      797    26.56    0.89  0.639
Residuals 15892 474127    29.83
[1] "year"
      Df Sum Sq Mean Sq F value Pr(>F)
train[, i]      6  10776  1796.0   61.59 <2e-16 ***
Residuals 15916 464147    29.2
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
[1] "month"
      Df Sum Sq Mean Sq F value Pr(>F)
train[, i]     11   1625   147.70    4.965 9.41e-08 ***
Residuals 15911 473299    29.75
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
[1] "day"
      Df Sum Sq Mean Sq F value Pr(>F)
train[, i]      6     299    49.80    1.67  0.124
Residuals 15916 474625    29.82
[1] "hour"
      Df Sum Sq Mean Sq F value Pr(>F)
train[, i]     23   3176   138.10   4.654 9.38e-13 ***
Residuals 15899 471747    29.67
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

After running ANOVA(Analysis of Variance) for the categorical variables, we can see p values as less than 0.05 for day and the date variables. We can remove these variables.





Fig. 2.3 : Correlation plot

Modelling

3.1 Model Selection

Model Selection is where we select the suitable modelling technique based on the type of dependent variable. Since the cab fare is a continuous variable we are going with Linear Regression and Random Forest(Classification) technique.

3.2 Multiple Linear Regression

In linear regression, the relationships are modelled using linear predictor functions whose unknown model parameters are estimated from the data. Such models are called linear models.

We can see that the Adjusted R squared is 54.6% which means we can explain 54.6% of the data using our model. F-statistic is 334 and p-value is $2.2e-16$ which can reject the null hypothesis that target variable does not depend on any of the predictor variables.

The MAPE is 38.7% and F-statistic is 334. Hence the accuracy of the model is 61.3%. This means our model is not good.

```
Residual standard error: 3.686 on 12691 degrees of freedom  
Multiple R-squared:  0.548,    Adjusted R-squared:  0.5463  
F-statistic: 334.5 on 46 and 12691 DF,  p-value: < 2.2e-16
```

3.3 Random Forest

Apart from Regression, we shall use one Classification model for the predictions. Number of trees used in this case is 500. MAPE for the model is 20.55%. Hence the accuracy is 79.45%.

Conclusion

4.1 Model Evaluation

We have 2 models for predicting the cab fare amount. Now, we need to decide on which one to choose.

We can compare the models using any of the following criteria :

1. Predictive Performance
2. Interpretability
3. Computational Efficiency

For our case, criteria 2 and 3 does not hold good. Hence we go with Predictive Performance.

Predictive performance can be measured by comparing Predictions of the models with real values of the target variables, and calculating some average error measure.

4.2 Error Metrics

1) Mean Absolute Percentage Error (MAPE):

MAPE measures the average magnitude of errors in a set of predictions without considering directions. It's an percentage average of difference between predicted and actual values .

2) Root Mean Square Error (RMSE):

RMSE is more of a magnitude of error calculation metrics. It's the square root of the average of squared difference between predicted and actual values.

Even though MAPE and RMSE both express average model prediction errors, RMSE is squared before it is averaged, it gives more weight to large errors compared to MAPE. Since these errors are negatively oriented(lower the better), it's good to consider MAPE for our error metrics.

4.3 Mean Absolute Percentage Error(MAPE)

MAPE can be calculated by the below formula and MAPE for Linear Regression is 38.7% and For Random Forest it is 20.55%.

$$\text{MAPE} = ((\text{Actual_value} - \text{Predicted Value})/\text{Actual Value})*100$$

4.3 Model Selection

As we have seen that from the below table, Random Forest model has better Accuracy and less MAPE. So we can choose Random Forest

Model	MAPE	Accuracy
Linear Regression - Model1	38.7%	61.3%
Random Forest - Model 2	20.55%	79.45%

Table 4.1 Model Accuracy

Chapter 5

R- Code

```
##### CAB FARE PREDICTION
```

```
#####
```

```
#Clean the environment
```

```
rm(list = ls(all=T))
```

```
#Set Working Directory
```

```
setwd("C:/Users/shrid/Downloads/edWisor/Project 2")
```

```
getwd()
```

```
#Load Libraries
```

```
library("readr")
```

```
library("plyr")
```

```
library("dplyr")
```

```
library("DMwR")
```

```
library("ggplot2")
```

```
library("corrgram")
```

```
library("rlist")
```

```
library("purrr")
```

```
library("geosphere")
```

```
library("car")
```

```
library("randomForest")
```


###LOAD THE DATA###

#Load the training data and have a glimpse of the data

```
train = read.csv(file = "train_cab.csv", header = TRUE, sep = ",", na.strings = c("", " ", "NA"))
```

```
head(train)
```

```
class(train)
```

```
names(train)
```

```
str(train) # Checking for structure of the data
```

```
summary(train)
```

#Convert variables into appropriate datatypes

```
train$fare_amount <- as.numeric(as.character(train$fare_amount))
```

```
typeof(train$fare_amount)
```

```
summary(train)
```

###DATA PRE-PROCESSING###

#Missing value Ananlysis

#Let's check for missing values

```
missing_values = sapply(train, function(x){sum(is.na(x))})
```

```
missing_values
```

#We have 25 missing values in fare_amount and 55 in passenger_count

#Let's check it in terms of percentage

```
missing_values_percentage = (missing_values*100/nrow(train))
```

```
missing_values_percentage
```

#less than 30% so no need to delete any dependent variables. Let's just delete rows for simplicity

#Let's drop the rows with missing values for simplicity

```
train <- na.omit(train)
```

```
train
```

```
str(train)
```

```
missing_values = sapply(train, function(x){sum(is.na(x))})
```

```
missing_values
```

#We do not see any missing values after clean up

#Feature Engineering

#Let's draw features out of datetime stamp

```
train$pickup_datetime <- gsub(' UTC', '', train$pickup_datetime)
```

```
train$pickup_datetime
```

```
train$date <- as.Date(train$pickup_datetime)
```

```
train$date
```

```
train$year <- substr(train$pickup_datetime,1,4)
```

```
train$year
```

```
train$month <- substr(train$pickup_datetime,6,7)
```

```
train$month
```

```
train$day <- weekdays(as.POSIXct(train$date), abbreviate = F)
```

```
train$day
```

```
train$date <- substr(train$date, 9, 10)
```

```
train$date
```

```
train$hour <- substr(train$pickup_datetime, 12, 13)
```

```
train$hour
```

```
class(train$month)
```

#Let's check for missing values after Feature Engineering

```
missing_values = sapply(train, function(x){sum(is.na(x))})
```

```
missing_values
```

#There is one Missing row, let's remove that

```
train = na.omit(train)
```

```
missing_values = sapply(train, function(x){sum(is.na(x))})
```

missing_values

#Outlier Analysis

#Outlier Ananlysis for continous variables

```
numeric_data <- subset(train, select =  
c(fare_amount,pickup_longitude,pickup_latitude,dropoff_longitude,dropoff_latitude,passenge  
r_count))
```

```
cnames <- colnames(numeric_data)
```

```
cnames
```

```
summary(numeric_data)
```

#Boxplot of all numeric variables

```
for(i in 1:length(cnames))
```

```
{
```

```
  assign(paste0("b",i), ggplot(aes_string(y = cnames[i]), data = train) +
```

```
    stat_boxplot(geom = "errorbar", width = 0.5) +
```

```
    geom_boxplot(outlier.color = "red", fill = "green", outlier.shape = 16, outlier.size = 1,  
notch = FALSE)+
```

```
    theme(legend.position = "bottom") + labs(cnames[i])+
```

```
    ggtitle(paste("Box plot for", cnames[i]))
```

```
  )
```

```
}
```

```
gridExtra::grid.arrange(b1,b2,b3,b4,b5,b6,ncol=2)
```

#Let's treat the Outliers

#Fare amount

```
summary(train$fare_amount)
```

```
Q1 <- quantile(train$fare_amount, 0.25)
```

Q1 #Quantile 1

```
Q3 <- quantile(train$fare_amount, 0.75)
```

Q3 #Quantile 3

```
UL <- Q3 + (1.5*IQR(train$fare_amount))
```

UL #Upper Limit

```
LL <- Q1 - (1.5*IQR(train$fare_amount))
```

LL #Lower limit

#Let's replace the values outside of range with Lower Limit and Upper Limit

```
train[train$fare_amount < LL, "fare_amount"] <- LL
```

```
train[train$fare_amount > UL, "fare_amount"] <- UL
```

```
summary(train$fare_amount)
```

#Since the fare can not be in negative and zero, let's replace those values with NA

```
train$fare_amount[train$fare_amount <= 0.01] <- NA
```

#Let's check the NA values and remove them

```
sum(is.na(train$fare_amount))
```

```
train <- na.omit(train)
```

```
#Pick up longitude
```

```
summary(train$pickup_longitude)
```

```
Q1 <- quantile(train$pickup_longitude, 0.25)
```

```
Q1
```

```
Q3 <- quantile(train$pickup_longitude, 0.75)
```

```
Q3
```

```
UL <- Q3 + (1.5*IQR(train$pickup_longitude))
```

```
UL
```

```
LL <- Q1 - (1.5*IQR(train$pickup_longitude))
```

```
LL
```

```
train[train$pickup_longitude < LL, "pickup_longitude"] <- LL
```

```
train[train$pickup_longitude > UL, "pickup_longitude"] <- UL
```

```
summary(train$pickup_longitude)
```

```
#Pick up latitude
```

```
summary(train$pickup_latitude)
```

```
Q1 <- quantile(train$pickup_latitude, 0.25)
```

```
Q1
```

```
Q3 <- quantile(train$pickup_latitude, 0.75)
```

```
Q3
```

```
UL <- Q3 + (1.5*IQR(train$pickup_latitude))
```

```
UL
```

```
LL <- Q1 - (1.5*IQR(train$pickup_latitude))
```

```
LL
```

```
train[train$pickup_latitude < LL, "pickup_latitude"] <- LL
```

```
train[train$pickup_latitude > UL, "pickup_latitude"] <- UL
```

```
summary(train$pickup_latitude)
```

```
#Drop off Longitude
```

```
summary(train$dropoff_longitude)
```

```
Q1 <- quantile(train$dropoff_longitude, 0.25)
```

```
Q1
```

```
Q3 <- quantile(train$dropoff_longitude, 0.75)
```

```
Q3
```

```
UL <- Q3 + (1.5*IQR(train$dropoff_longitude))
```

```
UL
```

```
LL <- Q1 - (1.5*IQR(train$dropoff_longitude))
```

```
LL
```

```
train[train$dropoff_longitude < LL, "dropoff_longitude"] <- LL
```

```
train[train$dropoff_longitude > UL, "dropoff_longitude"] <- UL
```

```
summary(train$dropoff_longitude)
```

#Drop off latitude

```
summary(train$dropoff_latitude)
```

```
Q1 <- quantile(train$dropoff_latitude, 0.25)
```

```
Q1
```

```
Q3 <- quantile(train$dropoff_latitude, 0.75)
```

```
Q3
```

```
UL <- Q3 + (1.5*IQR(train$dropoff_latitude))
```

```
UL
```

```
LL <- Q1 - (1.5*IQR(train$dropoff_latitude))
```

```
LL
```

```
train[train$dropoff_latitude < LL, "dropoff_latitude"] <- LL
```

```
train[train$dropoff_latitude > UL, "dropoff_latitude"] <- UL
```

```
summary(train$dropoff_latitude)
```

#Passenger count

```
summary(train$passenger_count)
```

```
Q1 <- quantile(train$passenger_count, 0.25)
```

```
Q1
```



```
Q3 <- quantile(train$passenger_count, 0.75)
```

```
Q3
```

```
UL <- Q3 + (1.5*IQR(train$passenger_count))
```

```
UL
```

```
LL <- Q1 - (1.5*IQR(train$passenger_count))
```

```
LL
```

```
train[train$passenger_count < 0, "passenger_count"] <- 0
```

```
train[train$passenger_count > 6, "passenger_count"] <- 6 #As there can be a max of 6  
passengers
```

#Since there should be at least 1 passenger for the trip to count, let's replace all the values less than 1 with NA

```
train$passenger_count[train$passenger_count < 1] <- NA
```

```
sum(is.na(train$passenger_count))
```

#Let's remove the missing values

```
train = na.omit(train)
```

```
sum(is.na(train$passenger_count))
```

#Let's visualize the boxplot after removing outliers

```
for(i in 1:length(cnames))
```

```
{
```

```

assign(paste0("b",i), ggplot(aes_string(y = cnames[i]), data = train) +
  stat_boxplot(geom = "errorbar", width = 0.5) +
  geom_boxplot(outlier.color = "red", fill = "grey", outlier.shape = 16, outlier.size = 1,
notch = FALSE)+
  theme(legend.position = "bottom") + labs(y=cnames[i])+
  ggtitle(paste("Box plot for", cnames[i]))
)
}

```

```

gridExtra::grid.arrange(b1,b2,b3,b4,b5,b6,ncol=3)

```

#Let's calculate distance between 2 geo codes

```

distance1 = function(long1, lat1, long2, lat2){
  loadNamespace("purrr")
  loadNamespace("geosphere")

  l1 = purrr::map2(long1, lat1, function(x,y) c(x,y))
  l2 = purrr::map2(long2, lat2, function(x,y) c(x,y))

  dist = purrr::map2(l1, l2, function(x,y) geosphere::distHaversine(x, y))
  distance_m = list.extract(dist, position = 1)

  distance = distance_m / 1000.0 ;#in km

```

```

distance
}

for( i in (1:nrow(train)))
{
  train$distance[i] = distance1(train$pickup_longitude[i], train$pickup_latitude[i],
train$dropoff_longitude[i], train$dropoff_latitude[i])
}

head(train)

str(train)


# Outliers for Distance

bp_dist = ggplot(aes_string( x = "distance", y = "fare_amount"), data = train) +
  stat_boxplot(geom = "errorbar", width = 0.25) +
  geom_boxplot(outlier.color = "red", fill = "green", outlier.shape = 16, outlier.size = 1, notch
= FALSE) +
  ggtitle(paste("Boxplot for Distance with Outliers"))

bp_dist

summary(train$distance)

train$distance[train$distance < 1] <- mean(train$distance)

```

```
summary(train$distance)
```

```
# Exploratory data analysis
```

```
# Frequency distribution of Numeric variables
```

```
#Fare amount
```

```
a1 <- ggplot(data = train, aes( x = train$fare_amount)) + geom_histogram(fill = "royalblue",  
bins = 20) +
```

```
labs(x = "Fare Amount", y = "Frequency") + ggtitle("Distribution of Fare Amount")
```

```
a1
```

```
#Pick up Longitude
```

```
a2 <- ggplot(data = train, aes( x = train$pickup_longitude)) + geom_histogram( fill =  
"lightgreen") +
```

```
labs(x = "Pick up Longitude", y = "Frequency") + ggtitle("Distribution of Pick Up  
Longitude")
```

```
a2
```

```
#Pick up Latitude
```

```
a3 <- ggplot(data = train, aes( x = train$pickup_latitude)) + geom_histogram(fill = "green") +  
labs(x = "Pick up Latitude", y = "Frequency") + ggtitle("Distribution of Pick Up Latitude")
```

```
a3
```

```
#Drop off Longitude
```

```
a4 <- ggplot(data = train, aes( x = train$dropoff_longitude)) + geom_histogram( fill =
"springgreen") +
```

```
labs(x = "Drop off Longitude", y = "Frequency") + ggtitle("Distribution of Drop off
Longitude")
```

```
a4
```

#Drop off Latitude

```
a5 <- ggplot(data = train, aes( x = train$dropoff_latitude)) + geom_histogram( fill =
"limegreen") +
```

```
labs(x = "Drop off Latitude", y = "Frequency") + ggtitle("Distribution of Drop off Latitude")
```

```
a5
```

#Passenger Count

```
a6 <- ggplot(data = train, aes( x = train$passenger_count)) + geom_histogram(binwidth = 1,
fill = "powderblue") +
```

```
labs(x = "Passenger count", y = "Frequency") + ggtitle("Distribution of Passenger count")
```

```
a6
```

#Distance

```
a7 <- ggplot(data = train, aes( x = train$distance)) + geom_histogram(binwidth = 1, fill =
"navy") +
```

```
labs(x = "Distance", y = "Frequency") + ggtitle("Distribution of Distance")
```

a7 #distance is right skewed. Let's remove skewness by applying log

```
logof10 = function(x) {
```

```
  ifelse(abs(x) <= 1, 0, sign(x)*log10(abs(x)))
```

```
}
```

```
train$distance = logof10(train$distance)
```

#Distance after Log transformation

```
a7 <- ggplot(data = train, aes( x = train$distance)) + geom_histogram(binwidth = 0.1, fill =  
"navy") +
```

```
labs(x = "Distance", y = "Frequency") + ggtitle("Distribution of Distance after Log  
Transformation")
```

```
a7
```

#Categorical variables

#Fare amount vs Year

```
b1 <- ggplot(train, aes(train$year, train$fare_amount)) + geom_bar(stat = "identity", fill =  
"yellow") +
```

```
labs( x = "Year", y = "Fare Amount") + ggtitle("Fare Amount vs Year")
```

```
b1
```

#Fare Amount vs Month

```
b2 <- ggplot(train, aes(train$month, train$fare_amount)) + geom_bar(stat = "identity", fill =  
"khaki") +
```

```
labs( x = "Month", y = "Fare Amount") + ggtitle("Fare Amount vs Month")
```

```
b2
```

#Fare Amount vs Day

```
b3 <- ggplot(train, aes(train$day, train$fare_amount)) + geom_bar(stat = "identity", fill = "orange") +
```

```
labs(x = "Day", y = "Fare Amount") + ggtitle("Fare Amount vs Day")
```

```
b3
```

```
#Fare Amount vs Hour
```

```
b4 <- ggplot(train, aes(train$hour, train$fare_amount)) + geom_bar(stat = "identity", fill = "peru") +
```

```
labs(x = "Hour", y = "Fare Amount") + ggtitle("Fare Amount vs Hour")
```

```
b4
```

```
#Correlation plot
```

```
c1 <- corrgram(train, order = FALSE, main = "Correlation plot")
```

```
c1
```

```
c2 <- corrgram(train, order = FALSE, main = "Correlation plot", panel = panel.cor)
```

```
c2
```

```
#ANOVA for categorical variables
```

```
categorical <- c("date", "year", "month", "day", "hour")
```

```

for(i in categorical){

  print(i)

  anova = summary(aov(fare_amount~train[,i],train))

  print(anova)

}

#p > 0.05 for day and date. So let's remove those variables

names(train)

data1 <- subset(train, select =
c(fare_amount,pickup_longitude,pickup_latitude,dropoff_longitude,dropoff_latitude,passenge
r_count,year,month,hour,distance))

data1

#Let's create train and test dataset

index = sample(1:nrow(data1), as.integer(0.8*nrow(data1)) )

train_data = data1[index,]

test_data = data1[-index,]

#Linear regression

model1 = lm(fare_amount~., train_data)

summary(model1)

par(mfrow = c(2,2))

```



```
plot(model1)
```

```
vif(model1) #less than 5
```

```
dwt(model1)
```

```
prediction1 <- predict(model1, test_data[,-1])
```

```
df = data.frame("actual" = test_data[,1], "pred" = prediction1)
```

```
df
```

```
MAPE = function(y, yhat){
```

```
  mean(abs((y-yhat)/y))
```

```
}
```

```
MAPE(test_data[,1],prediction1)
```

```
#Random Forest
```

```
model2 <- randomForest(fare_amount~., train_data, tree = 500)
```

```
summary(model2)
```

```
prediction2 <- predict(model2, test_data[,-1])
```

```
df = cbind(df, prediction2)
```

```
df
```

```
MAPE(test_data[,1],prediction2)
```

```
#Model Evaluation
```

```
test <- read.csv(file = "test.csv", header = TRUE, sep = ",", na.strings = c("", " ", "NA"))
```

```
head(test)
```

```
str(test)
```

```
summary(test)
```

```
#Feature engineering
```

```
test$pickup_datetime <- gsub(" UTC", "", test$pickup_datetime)
```

```
test$date <- as.Date(test$pickup_datetime)
```

```
test$year <- substr(as.character(test$pickup_datetime),1,4)
```

```
test$month <- substr(as.character(test$pickup_datetime),6,7)
```

```
test$day <- weekdays.POSIXt(test$date, abbreviate = FALSE)
```

```
test$date <- substr(as.character(test$date),9,10)
```

```
test$hour <- substr(as.factor(test$pickup_datetime), 12, 13)
```

```
str(test)
```

```
# Data preprocessing
```

```
#Missing value
```

```
sum(is.na(test))
```

```
#Outlier
```

```
summary(test)
```

```
#distance
```

```
for( i in (1:nrow(test)))
```

```
{
```

```
  test$distance[i] = distance1(test$pickup_longitude[i], test$pickup_latitude[i],  
test$dropoff_longitude[i], test$dropoff_latitude[i])
```

```
}
```

```
head(test)
```

```
summary(test$distance)
```

```
test$distance[test$distance < 1] = mean(test$distance)
```

```
#Distribution of distance
```

```
z <- ggplot(test, aes_string(x = test$distance)) + geom_histogram() + geom_density()
```

```
z
```

```
test$distance <- logof10(test$distance)
```

```
summary(test$distance)
```

```
#Modelling
```

```
Final <- subset(test, select =  
c(pickup_longitude,pickup_latitude,dropoff_longitude,dropoff_latitude,passenger_count,year,  
month,hour,distance ))
```

```
Final_model <- randomForest(fare_amount~. , data1, ntree = 500, method = "anova")
```

```
predictions <- predict(Final_model, Final)
```

```
predictions
```

```
test$fare_amount <- predictions
```

```
summary(test)
```

```
summary(test$fare_amount)
```

```
head(test)
```

Chapter 6

Python Code

```
#####CAB FARE PREDICTION#####  
#####
```

```
#Import the library
```

```
import pandas as pd
```

```
import os
```

```
import numpy as np
```

```
from ggplot import *
```

```
import seaborn as sns
```

```
#set the Working Directory
```

```
os.chdir("C:¥¥Users¥shrid¥Downloads¥edWisor¥Project 2")
```

```
os.getcwd()
```

```
#Load the data
```

```
train = pd.read_csv("train_cab.csv")
```

```
#Have a glimpse of the data
```

```
train.head()
```

```
train.shape
```

```
#Data type check
```

```
train.dtypes
```

```
#Summary of the data
```

```
train.describe()
```

```

# Let's convert fare_amount to numeric

train['fare_amount'] = pd.to_numeric(train['fare_amount'], errors = 'coerce')

train['fare_amount'].dtypes

train.describe()

#train['pickup_datetime'] = pd.to_datetime(train['pickup_datetime'], format='%Y-%m-%d
%H:%M:%S UTC')

#Getting Error because of some value, let's make it NA and drop

train.loc[train['pickup_datetime'] == '43', 'pickup_datetime'] = np.nan

train = train.drop(train[train['pickup_datetime'].isnull()].index, axis = 0)

train['pickup_datetime'] = pd.to_datetime(train['pickup_datetime'], format='%Y-%m-%d
%H:%M:%S UTC')


#Checking for data types and top 5 rows

train.dtypes

train.head()


#Let's check for Missing values

train.isnull().sum()


#Let's remove Missing values

train = train.drop(train[train['fare_amount'].isnull()].index,axis=0)

train = train.drop(train[train['passenger_count'].isnull()].index,axis=0)

#Let's check if there are any missing values after removal

train.isnull().sum()

train.shape

```

#Feature Engineering

#Let's extract Date, year, month, day, hour out of the timestamp

```
train['year'] = train['pickup_datetime'].dt.year
train['month'] = train['pickup_datetime'].dt.month
train['date'] = train['pickup_datetime'].dt.day
train['day'] = train['pickup_datetime'].dt.dayofweek
train['hour'] = train['pickup_datetime'].dt.hour
```

#Let's have a glimpse of the data after extracting new variables

```
train.shape
train.head()
```

#Let's drop the pickup_datetime variable

```
train = train.drop('pickup_datetime',axis=1)
train.head()
```

#Outlier analysis

```
cnames = ['fare_amount', 'pickup_longitude', 'pickup_latitude',
          'dropoff_longitude', 'dropoff_latitude', 'passenger_count']
import matplotlib.pyplot as plt
%matplotlib inline
for i in cnames:
    print(i)
    plt.boxplot(train[i])
    plt.xlabel(i)
```

```
plt.ylabel('fare_amount')
```

```
plt.title('outlier analysis')
```

```
plt.show()
```

#Let's treat outliers

#Fare Amount

Quartiles

```
Q1,Q3 = np.percentile(train['fare_amount'],[25,75])
```

#IQR

```
IQR = Q3-Q1
```

Lower and upper limits

```
LL = Q1 - (1.5 * IQR)
```

```
UL = Q3 + (1.5 * IQR)
```

Capping with ul for maximum values

```
train.loc[train['fare_amount'] < LL , 'fare_amount'] = LL
```

```
train.loc[train['fare_amount'] > UL , 'fare_amount'] = UL
```

#Less than 0 and 0.01 values

```
train.loc[train.fare_amount <= 0, 'fare_amount'] = np.nan
```

```
train.loc[train.fare_amount == 0.01, 'fare_amount'] = np.nan
```

```
train = train.drop(train[train['fare_amount'].isnull()].index, axis = 0)
```

```
train['fare_amount'].describe()
```


#Pick up Longitude

```
train['pickup_longitude'].describe()
```

Quartiles

```
Q1,Q3 = np.percentile(train['pickup_longitude'],[25,75])
```

#IQR

```
IQR = Q3-Q1
```

Lower and upper limits

```
LL = Q1 - (1.5 * IQR)
```

```
UL = Q3 + (1.5 * IQR)
```

Max of this variable is 40.77 which we can consider as outlier and capping with UL

```
train.loc[train['pickup_longitude'] < LL , 'pickup_longitude'] = LL
```

```
train.loc[train['pickup_longitude'] > UL , 'pickup_longitude'] = UL
```

#Pick up Latitude

```
train['pickup_latitude'].describe()
```

Quartiles

```
Q1,Q3 = np.percentile(train['pickup_latitude'],[25,75])
```

#IQR

```
IQR = Q3-Q1
```

Lower and upper limits

```
LL = Q1 - (1.5 * IQR)
```

```
UL = Q3 + (1.5 * IQR)
```

Capping with ul for maxmimu values

```
train.loc[train['pickup_latitude'] < LL , 'pickup_latitude'] = LL
```

```
train.loc[train['pickup_latitude'] > UL , 'pickup_latitude'] = UL
```

```
train['pickup_latitude'].describe()
```

```
#Drop off Longitude
```

```
train['dropoff_longitude'].describe()
```

```
# Quartiles
```

```
Q1,Q3 = np.percentile(train['dropoff_longitude'],[25,75])
```

```
#IQR
```

```
IQR = Q3-Q1
```

```
# Lower and upper limits
```

```
LL = Q1 - (1.5 * IQR)
```

```
UL = Q3 + (1.5 * IQR)
```

```
# Capping with ul for maximum values
```

```
train.loc[train['dropoff_longitude'] < LL,'dropoff_longitude'] = LL
```

```
train.loc[train['dropoff_longitude'] > UL,'dropoff_longitude'] = UL
```

```
#Drop off Latitude
```

```
train['dropoff_latitude'].describe()
```

```
# Quartiles
```

```
Q1,Q3 = np.percentile(train['dropoff_latitude'],[25,75])
```

```
#IQR
```

```
IQR = Q3-Q1
```

```
# Lower and upper limits
```

```
LL = Q1 - (1.5 * IQR)
```

```
UL = Q3 + (1.5 * IQR)
```

Capping with ul for maximum values

```
train.loc[train['dropoff_latitude'] < LL, 'dropoff_latitude'] = LL
train.loc[train['dropoff_latitude'] > UL, 'dropoff_latitude'] = UL
train['dropoff_latitude'].describe()
```

#Passenger Count

```
train['passenger_count'].describe()
```

Quartiles

```
Q1, Q3 = np.percentile(train['passenger_count'], [25, 75])
```

#IQR

```
IQR = Q3 - Q1
```

Lower and upper limits

```
LL = round(Q1 - (1.5 * IQR))
```

```
UL = round(Q3 + (1.5 * IQR))
```

Capping with UL for maximum values

```
train.loc[train['passenger_count'] < LL, 'passenger_count'] = LL
```

```
train.loc[train['passenger_count'] > 6, 'passenger_count'] = UL
```

```
train.loc[train['passenger_count'] < 1, 'passenger_count'] = np.nan
```

```
train = train.drop(train[train['passenger_count'].isnull()].index, axis = 0)
```

```
train['passenger_count'].describe()
```

#Boxplot after removal of Outliers

```
import matplotlib.pyplot as plt
```

```
%matplotlib inline
```

```
for i in cnames:
```

```
    print(i)
```

```
    plt.boxplot(train[i])
```

```
    plt.xlabel(i)
```

```
    plt.ylabel('fare_amount')
```

```
    plt.title('outlier analysis')
```

```
    plt.show()
```

```
train.head()
```

```
train.shape
```

#Let's calculate distance between two geo codes using haversine formula

```
from math import radians, cos, sin, asin, sqrt
```

```
def distance1(lat1, long1, lat2, long2):
```

```
    R_earth = 6371 # earth radius (km)
```

```
    #Convert degrees to radians
```

```
    lat1, long1, lat2, long2 = map(np.radians, [lat1, long1, lat2, long2])
```

```
    #Compute distances along lat, lon dimensions
```

```
    dlat = lat2 - lat1
```

```
    dlon = long2 - long1
```

#Compute haversine distance

```
a = np.sin(dlat/2.0)**2 + np.cos(lat1) * np.cos(lat2) * np.sin(dlon/2.0)**2
```

```
return 2 * R_earth * np.arcsin(np.sqrt(a))
```

#Let's calculate distance using above function

```
train['distance'] =  
distance1(train['pickup_latitude'],train['pickup_longitude'],train['dropoff_latitude']  
,train['dropoff_longitude'])
```

```
train.head()
```

```
train.shape
```

```
train.dtypes
```

#Outlier for distance

```
%matplotlib inline
```

```
plt.boxplot(train['distance'])
```

```
plt.xlabel('distance')
```

```
plt.title('outlier analysis')
```

```
plt.show()
```

```
train['distance'].describe()
```

#Let's replace less than 1 values with mean

```
train.loc[train.distance < 1,'distance'] = train['distance'].mean()
```

#Boxplot after removing outliers

```
%matplotlib inline
```

```
plt.boxplot(train['distance'])
```

```
plt.xlabel('distance')
```

```
plt.title('outlier analysis')
```

```
plt.show()
```

```
train.head()
```

```
train.shape
```

#Exploratory Data Analysis

#Continuous variables analysis

#Fare amount

```
a1 = ggplot(train, aes( x = 'fare_amount')) + geom_histogram(fill = "royalblue", bins = 25) + ¥
```

```
labs(x = "Fare Amount", y = "Frequency") + ggtitle("Distribution of Fare Amount")
```

```
a1
```

#Pickup Longitude

```
a2 = ggplot(train, aes( x = 'pickup_longitude')) + geom_histogram( fill = "lightgreen") + ¥
```

```
labs(x = "Pick up Longitude", y = "Frequency") + ggtitle("Distribution of Pick Up  
Longitude")
```

```
a2
```

#Pickup Latitude

```
a3 = ggplot(train, aes( x = 'pickup_latitude')) + geom_histogram(fill = "green") + ¥
```

```
labs(x = "Pick up Latitude", y = "Frequency") + ggtitle("Distribution of Pick Up Latitude")
```

a3

#Drop off Longitude

```
a4 = ggplot(train, aes( x = 'dropoff_longitude')) + geom_histogram( fill = "springgreen") +  
  labs(x = "Drop off Longitude", y = "Frequency") + ggtitle("Distribution of Drop off  
Longitude")
```

a4

#Drop off Latitude

```
a5 = ggplot(train, aes( x = 'dropoff_latitude')) + geom_histogram( fill = "limegreen") +  
  labs(x = "Drop off Latitude", y = "Frequency") + ggtitle("Distribution of Drop off Latitude")
```

a5

#Passenger Count

```
a6 = ggplot(train, aes( x = 'passenger_count')) + geom_histogram(binwidth = 1, fill =  
"powderblue") +  
  labs(x = "Passenger count", y = "Frequency") + ggtitle("Distribution of Passenger count")
```

a6

#Distance

```
a7 = ggplot(train, aes( x = 'distance')) + geom_histogram(binwidth = 1, fill = "navy") +  
  labs(x = "Distance", y = "Frequency") + ggtitle("Distribution of Distance")
```

a7

#Distance is right skewed, let's remove skewness using log transformation

```
train['distance'] = np.log(train['distance'])
```

#Let's check distribution after transformation

```
a7 = ggplot(train, aes( x = 'distance')) + geom_histogram( fill = "navy") +  
  labs(x = "Distance", y = "Frequency") + ggtitle("Distribution of Distance")
```

a7

#Categorical Variables

#Fare amount vs Year

```
b1 = ggplot(train, aes('year', 'fare_amount')) + geom_bar(stat = "identity", fill = "yellow") +  
  labs(x = "Year", y = "Fare Amount") + ggtitle("Fare Amount vs Year")  
b1
```

#Fare amount vs Month

```
b2 = ggplot(train, aes('month', 'fare_amount')) + geom_bar(stat = "identity", fill = "khaki") +  
  labs(x = "Month", y = "Fare Amount") + ggtitle("Fare Amount vs Month")  
b2
```

#Fare amount vs Day

```
b3 = ggplot(train, aes('day', 'fare_amount')) + geom_bar(stat = "identity", fill = "orange") +  
  labs(x = "Day", y = "Fare Amount") + ggtitle("Fare Amount vs Day")  
b3
```

#Fare amount vs Hour

```
b4 = ggplot(train, aes('hour', 'fare_amount')) + geom_bar(stat = "identity", fill = "peru") +  
  labs(x = "Hour", y = "Fare Amount") + ggtitle("Fare Amount vs Hour")  
b4
```

#Correlation plot

```
cnames = ['fare_amount', 'pickup_longitude', 'pickup_latitude',  
          'dropoff_longitude', 'dropoff_latitude', 'passenger_count', 'distance']  
  
df_corr = train.loc[:, cnames]  
c1 = df_corr.corr()  
print(c1)
```



```

sns.heatmap(c1,square=True,annot=True)

#Anova test for categorical variables

cat = ['year','month', 'date', 'day', 'hour']


import statsmodels.api as sm


from statsmodels.formula.api import ols


for i in cat:

    mod = ols('fare_amount' + '~' + i, data = train).fit()

    aov_table = sm.stats.anova_lm(mod, typ = 2)

    print(aov_table)

#p > 0.05 for day and date. So let's remove those variables

DropVar = ['date', 'day']


train = train.drop(DropVar , axis = 1)


#Let's create Train and Test dataset

from sklearn.model_selection import train_test_split


train_data,test_data = train_test_split(train, test_size = 0.2, random_state = 123)

train_data.head()


#Linear regression

import statsmodels.api as sm


from sklearn.metrics import mean_squared_error

```

```

from sklearn.linear_model import LinearRegression

from sklearn.metrics import r2_score


#Let's train the model

model1 = sm.OLS(train_data.iloc[:,0].astype(float), train_data.iloc[:,1:10].astype(float)).fit()

model1.summary()


#Let's predict

prediction1 = model1.predict(test_data.iloc[:,1:10])

data_model1 = pd.DataFrame({'actual': test_data.iloc[:,0], 'pred': prediction1})

data_model1.head()


#Function to calculate MAPE

def MAPE(y_actual,y_pred):

    mape = np.mean(np.abs((y_actual - y_pred)/y_actual)*100)

    return mape

MAPE(test_data.iloc[:,0],prediction1)


#Random Forest

from sklearn.ensemble import RandomForestRegressor


#Let's train the model

model2 =

RandomForestRegressor(n_estimators=500,random_state=123).fit(train_data.iloc[:,1:10],
train_data.iloc[:,0])

```

#Let's predict

```
prediction2 = model2.predict(test_data.iloc[:,1:10])
```

```
data_model2 = pd.DataFrame({"actual" : test_data.iloc[:,0],"pred" : prediction2})
```

```
data_model2.head()
```

```
MAPE(test_data.iloc[:,0], prediction2)
```

#Model Evaluation

#Load the test data

```
test = pd.read_csv('test.csv')
```

#Glimpse of the data

```
test.head()
```

```
test.dtypes
```

```
test.describe()
```

#Format conversion

```
test['pickup_datetime'] = pd.to_datetime(test['pickup_datetime'], format='%Y-%m-%d  
%H:%M:%S UTC')
```

#Feature Engineering

```
test['year'] = test['pickup_datetime'].dt.year
```

```
test['month'] = test['pickup_datetime'].dt.month
```

```
test['date'] = test['pickup_datetime'].dt.day
```

```
test['day'] = test['pickup_datetime'].dt.dayofweek
```

```
test['hour'] = test['pickup_datetime'].dt.hour
```

#Let's delete Pick up datetime

```
test = test.drop(['pickup_datetime'], axis = 1)
```

```
test.head()
```

```
#Let's check for Missing values
```

```
test.isnull().sum()
```

```
#Outlier Analysis
```

```
test.describe()
```

```
#Let's calculate Distance
```

```
test['distance'] =
```

```
distance1(test['pickup_latitude'],test['pickup_longitude'],test['dropoff_latitude'],test['dropoff_
```

```
ongitude'])
```

```
test.head()
```

```
#Outliers for Distance
```

```
test['distance'].describe()
```

```
#Let's replace 0 with mean distance values
```

```
test.loc[test.distance < 1,'distance'] = test['distance'].mean()
```

```
test['distance'].describe()
```

```
test.head()
```

```
test = test.drop(['date', 'day'], axis = 1)
```

```
#Columns in Test data
```

```
test.columns
```

```
#Outliers in Distance
```

```
sns.distplot(test['distance'],color='black')
```

```
plt.title("Distribution of variable distance")
```

```
plt.ylabel("Density")
```

```
plt.show()
```

```
#Right skewed, Let's remove skewness using log function
```

```
test['distance'] = np.log(test['distance'])
```

```
#Data after log transformation
```

```
sns.distplot(test['distance'],color='black')
```

```
plt.title("Distribution of distance")
```

```
plt.ylabel("Density")
```

```
plt.show()
```

```
#Let's apply Random Forest model
```

```
Final_model = model2.predict(test)
```

```
#Dependent variable
```

```
test['Predicted_Fare_Amount'] = Final_model
```

```
test.head()
```